

8: Joint Random Variables / Multivariate Distributions

PSTAT 120A: Summer 2022

Ethan P. Marzban

July 12, 2022

University of California, Santa Barbara

- Axioms of Probability, Probability Spaces, Counting
- Conditional Probabilities, independence, etc.
- Basics of Random Variables (classification, p.m.f., c.m.f., moments)
- Discrete Distributions
- Continuous Distributions
- Transformations of Random Variables
- Double Integrals

Random Vectors

- Consider the following experiment: suppose I pick a point P at random from the interior of the unit disk, and I let X denote the x -coordinate and Y denote the y -coordinate.
- We could investigate X and Y separately, but we have this intuitive sense that these two random variables are in some way related.
- We will quantify this relationship in an upcoming lecture. For now, we will simply say: “let’s try and consider X and Y *together*, as a pair (X, Y) .”
- We can imagine generalizing this to not just two random variables, but a collection of n random variables! (X_1, X_2, \dots, X_n) .
- Now, remember how I said a random variable X maps from Ω to \mathbb{R} ? Well, clearly when we start to imagine pairs (or tuples) of random variables we no longer have a map from Ω to \mathbb{R} .
- Specifically, let’s consider that “picking a point” example; Ω is simply the unit disk $\Omega = \{(x, y) : x^2 + y^2 \leq 1\}$. Additionally, this pair (X, Y) takes an element in Ω and **spits out a pair of numbers** (namely, the x - and y -coordinates of the point, respectively). In other words,

$$(X, Y) : \Omega \rightarrow \mathbb{R}^2$$

- For this reason, we often refer to the pair (X, Y) as a **random vector** as opposed to a random variable. (Another terminology is to call them a **pair of bivariate random variables**, but this language does not generalize as nicely to more than 2

- Let's start making some of this more formal.

Definition: Random Vector

Given a probability space $(\Omega, \mathcal{F}, \mathbb{P})$, a **random vector**

$$\vec{X} = \begin{pmatrix} X_1 \\ X_2 \\ \vdots \\ X_n \end{pmatrix}$$

is a mapping $\vec{X} : \Omega \rightarrow \mathbb{R}^n$. We say that the **dimension** of \vec{X} is n , or that \vec{X} is an **n -dimensional** random vector.

- Though it is customary to write vectors in column format, often times we are lazy and simply write them as row vectors:

$$\vec{X} = (X_1, X_2, \dots, X_n)$$

- Remember how we constructed continuous random variables? Given a probability space $(\Omega, \mathcal{F}, \mathbb{P})$ and a [continuous] random variable $X : \Omega \rightarrow \mathbb{R}$, we argued that depending on our choice of \mathbb{P} we can construct a c.d.f. $F_X(x) := \mathbb{P}(X \leq x)$, which, provided we have differentiability, gave rise to a p.d.f. that we can use to find probabilities, expectations, etc.
- We can do something similar for random vectors. We start with the notion of a:

Definition: Joint Cumulative Distribution Function

Given an n -dimensional random vector $\vec{X} = (X_1, X_2, \dots, X_n)$ we define the **joint cumulative distribution function** (or **joint c.d.f.**, for short) to be

$$F_{X_1, X_2, \dots, X_n}(x_1, x_2, \dots, x_n) := \mathbb{P}(X_1 \leq x_1, X_2 \leq x_2, \dots, X_n \leq x_n)$$

- Under appropriate conditions, we have the following:

Theorem

Under certain conditions (conditions over which we won't concern ourselves for the purposes of this class), we have the existence of a function $f_{X_1, X_2, \dots, X_n}(x_1, x_2, \dots, x_n)$ such that

$$\begin{aligned} F_{X_1, X_2, \dots, X_n}(x_1, x_2, \dots, x_n) \\ = \int_{-\infty}^{x_n} \cdots \int_{-\infty}^{x_2} \int_{-\infty}^{x_1} f_{X_1, X_2, \dots, X_n}(t_1, t_2, \dots, t_n) dt_1 dt_2 \cdots dt_n \end{aligned}$$

Such a function is called a **joint probability density function** (a.k.a. **joint p.d.f**, or just **joint density**).

Theorem

A joint density function must satisfy the following two conditions:

(1) $f_{X_1, \dots, X_n}(x_1, \dots, x_n) \geq 0$ for all $(x_1, \dots, x_n) \in \mathbb{R}^n$

(2) $\int \cdots \int_{\mathbb{R}^n} f_{X_1, \dots, X_n}(x_1, \dots, x_n) \, dx_1 \cdots dx_n = 1$

This also works in the other direction; that is, if we have a function $f_{X_1, \dots, X_n}(x_1, \dots, x_n)$ that satisfies the above two conditions then it is the joint density of some random vector \vec{X} .

- This is perhaps a good time to introduce some simplifying notation.
- When dealing with random vectors in generality, we often will need to write n -dimensional integrals.
- I shall adopt the following notation, which I borrow from Physics:

$$\int_{\mathbb{R}^n} f_{\vec{X}}(\vec{x}) \, d\vec{x}$$

shall mean

$$\int \cdots \int_{\mathbb{R}^n} f_{X_1, \dots, X_n}(x_1, \dots, x_n) \, dx_1 \cdots dx_n$$

- So, for instance, the second condition above can be written as $\int_{\mathbb{R}^n} f_{\vec{X}}(\vec{x}) \, d\vec{x} = 1$.
- By the way: in the subscript I'm using a capital X (\vec{X}) and in the argument I'm using a lowercase x (\vec{x}).

- Okay, I admit that dealing with random vectors in generality can get a bit pesky.
- When you start talking about “sampling” in 120B, you’ll see why random vectors arise **extremely** often throughout statistics. (Loosely speaking: Statisticians like to collect a *lot* of data, which can be modeled nicely using random vectors; a random variable for each observation!)
- For the purposes of this class, we will primarily restrict our considerations to $n = 2$, which gives rise to so-called **bivariate** random variables and distributions. But let’s quickly run through some generalities first:

Multivariate distributions

- Much like we had distributions in the case of random variables, we also have distributions in the case of random vectors. These distributions are often referred to as **multivariate** distributions.
- Unlike with univariate distributions, however, there aren't a whole lot that have specific names associated with them.
- There are two exceptions; we will discuss one of them in a bit, and time permitting we will discuss the second a little later.
- Let's return to our "picking a point" example. More generally, we could consider the following situation: from a region \mathcal{R} in \mathbb{R}^n , pick a point at random.
- Associated with this experiment, we could utilize the following choice of probability measure:

$$\mathbb{P}(A) = \frac{\text{volume}(A)}{\text{volume}(\Omega)}$$

In the case of $n = 2$, this is equivalently written as

$$\mathbb{P}(A) = \frac{\text{area}(A)}{\text{area}(\Omega)}$$

- Letting $\vec{X} = (X_1, \dots, X_n)$ denote the coordinates of the selected points, one can find (through a similar argument we used to derive the p.d.f. of the $\text{Unif}[a, b]$ distribution) that the joint density of \vec{X} is

$$f_{\vec{X}}(\vec{x}) = \frac{1}{\text{area}(\Omega)} \cdot \mathbb{1}_{\{\vec{x} \in \Omega\}} \quad (1)$$

- So, for instance, in our “picking a point from the unit disc” problem the joint density of (X, Y) is

$$f_{X,Y}(x, y) = \frac{1}{\pi} \cdot \mathbb{1}_{\{(x,y): x^2+y^2 \leq 1\}} = \begin{cases} \frac{1}{\pi} & \text{if } x^2 + y^2 \leq 1 \\ 0 & \text{otherwise} \end{cases}$$

- You can check that this is in fact a valid joint probability density function!
- This distribution (i.e. the one with p.d.f. listed in equation (1) above) doesn't have a standard name, but I will often refer to this as a **multivariate uniform** distribution, due to its similarity to our familiar $\text{Unif}[a, b]$ distribution (note that an interval $[a, b]$ is nothing but a “region” in \mathbb{R}^1 !)

Bivariate Random Variables/Distributions

- Given a pair of random variables (X, Y) , we have the notion of a bivariate density function: a function $f_{X,Y}(x, y)$ that is nonnegative over \mathbb{R}^2 and also integrates to unity (when integrated over \mathbb{R}^2).
- With such a function, we find that a great many of our familiar functions have nice bivariate analogs: for example, the LOTUS becomes

$$\mathbb{E}[g(X, Y)] = \iint_{\mathbb{R}^2} g(x, y) \cdot f_{X,Y}(x, y) \, dA$$

- Additionally, just like we found probabilities in the univariate case by integrating the density, we get probabilities in the bivariate case by integrating the bivariate density:

$$\mathbb{P}((X, Y) \in \mathcal{R}) = \iint_{\mathcal{R}} f_{X,Y}(x, y) \, dA$$

- Maybe now you see why we did that whole double integral review...
- One new piece of terminology: the region over which a joint density is nonzero is called the **support** of the random vector. It will almost always be a good idea to sketch the support of a random vector!

- One more piece of terminology that is unique to random vectors is that of the **marginal density/distribution**:

Definition: Marginals

Given a random vector $\vec{X} = (X_1, \dots, X_n)$ with joint p.d.f. $f_{\vec{X}}(\vec{x})$, the **marginal density of X_i** is given by

$$\int_{-\infty}^{\infty} f_{\vec{X}}(\vec{x}) dx_i$$

In other words, to find the marginal density of X_i we simply integrate out X_i from the joint density.

Suppose (X, Y) is a pair of random variables with joint density given by

$$f_{X,Y}(x, y) = \begin{cases} c \cdot e^{-(x+y)} & \text{if } x \leq y < \infty, 0 \leq x < \infty \\ 0 & \text{otherwise} \end{cases}$$

where $c > 0$ is an as-of-yet undetermined constant.

- (a) Find the value of c that ensures $f_{X,Y}(x, y)$ is a valid joint p.d.f..
- (b) Compute $\mathbb{P}(X \geq 0.5, Y \geq 0.5)$
- (c) Compute $\mathbb{E}[XY]$
- (d) Find $f_X(x)$, the marginal density of X .

Discrete?

- So far we've dealt only with continuous random vectors. What about discrete ones?
- Well, the primary difference is that instead of a joint p.d.f. we have a (perhaps more easily intuited) joint probability *mass* function

$$p_{X_1, \dots, X_n}(x_1, \dots, x_n) = \mathbb{P}(X_1 = x_1, \dots, X_n = x_n)$$

that obeys:

- (1) $0 \leq p_{X_1, \dots, X_n}(x_1, \dots, x_n) \leq 1$ for all $\vec{x} \in \mathbb{R}^n$
 - (2) $\sum_{\mathbb{R}^n} p_{X_1, \dots, X_n}(x_1, \dots, x_n) = 1$
- Familiar analogies apply:

$$\mathbb{P}(\vec{X} \in A) = \sum_{\vec{x} \in A} p_{\vec{X}}(\vec{x})$$

and the LOTUS becomes

$$\mathbb{E}[g(\vec{X})] = \sum_{\mathbb{R}^n} g(\vec{x}) \cdot p_{\vec{x}}(\vec{x})$$

[note that both summations above are really n -summations; that is, they are n sums in one]

Let (X, Y) be a pair of bivariate discrete random variables with joint p.m.f.

$$p_{X,Y}(x, y) = \begin{cases} c \cdot xy & \text{if } x \in \{1, 2, 3, 4\}, y \in \{1, 2, 3\} \\ 0 & \text{otherwise} \end{cases}$$

where $c > 0$ is an as-of-yet undetermined constant.

- (a) Find the value of c
- (b) Compute $\mathbb{E}[XY]$

Theorem: Linearity of Expectation

Given a collection of random variables X_1, \dots, X_n and a collection of constants a_1, \dots, a_n , we have

$$\mathbb{E} \left[\sum_{i=1}^n a_i X_i \right] = \sum_{i=1}^n a_i \mathbb{E}[X_i]$$

Proof.

- We use the multidimensional LOTUS:

$$\begin{aligned} \mathbb{E} \left[\sum_{i=1}^n X_i \right] &= \int_{\mathbb{R}^n} \left(\sum_{i=1}^n a_i x_i \right) f_{\vec{X}}(\vec{x}) \, d\vec{x} \\ &= \int_{\mathbb{R}^n} \sum_{i=1}^n (a_i x_i f_{\vec{X}}(\vec{x})) \, d\vec{x} \\ &= \sum_{i=1}^n \left[\int_{\mathbb{R}^n} a_i x_i f_{\vec{X}}(\vec{x}) \, d\vec{x} \right] \\ &= \sum_{i=1}^n \mathbb{E}[a_i X_i] = \sum_{i=1}^n a_i \mathbb{E}[X_i] \end{aligned}$$

- If the vector notation on the previous slide is too confusing, you can think of things in terms of $n = 2$; the proof for general n follows analogously.

$$\begin{aligned}\mathbb{E}[a_1 X_1 + a_2 X_2] &= \iint_{\mathbb{R}^2} (a_1 x_1 + a_2 x_2) f_{X_1, X_2}(x_1, x_2) \, dA \\ &= \iint_{\mathbb{R}^2} [a_1 x_1 f_{X_1, X_2}(x_1, x_2) + a_2 x_2 a_1 x_1 f_{X_1, X_2}(x_1, x_2)] \, dA \\ &= \iint_{\mathbb{R}^2} a_1 x_1 f_{X_1, X_2}(x_1, x_2) \, dA + \iint_{\mathbb{R}^2} a_2 x_2 f_{X_1, X_2}(x_1, x_2) \, dA \\ &= \mathbb{E}[a_1 X_1] + \mathbb{E}[a_2 X_2] = a_1 \mathbb{E}[X_1] + a_2 \mathbb{E}[X_2]\end{aligned}$$